White Paper 6:

Assumptions, Implications, and the Case for Coherent Intelligence

Brayden McLean | 7Site LLC | Hot Springs, Arkansas

Theoretical framework: Celeste Sol Weaver

February 2026

Abstract. This paper serves three purposes. First, it enumerates every assumption underlying Trinity Infinity Geometry (TIG) in explicit, falsifiable form, ordered so that a peer reviewer can evaluate each independently. Second, it derives what follows if these assumptions hold, with honest separation between proven results, projected consequences, and speculative extensions. Third, it argues that TIG's results point toward a paradigm distinct from statistical machine learning: *coherent intelligence*, in which systems are organized by spectral classification of their dynamics rather than by pattern-matching over data. We do not claim TIG replaces artificial intelligence. We claim it identifies the mathematical structure that any intelligence — artificial or biological — must satisfy to remain stable, and that this structure is not captured by current AI architectures.

1. The Assumption Stack

TIG rests on twelve assumptions, organized from weakest (most likely to survive peer review) to strongest (most likely to be challenged). Each is stated as a falsifiable claim with its evidence status.

1.1 Layer 1: Established Mathematics (not ours to prove)

ID	Assumption	Statement	Status
A1	Quadratic iteration is universal.	Any smooth map can be locally approximated by a quadratic. The logistic [illegible]	Established. Textbook dynamical systems (Strog[illegible]
A2	The Koopman-von Neumann theorem holds.	Every classical dynamical system admits an exact Hilbert space formulation via the Koopman operator [illegible]	Established. Koopman (1931), Lasota-Mackey (19[illegible]
A3	Lyapunov exponents classify stability.	For a map with fixed point x*, the quantity Lambda = ln|f'(x*)| determines convergence (Lambda < 0), divergence (Lambda > [illegible]	Established. Oseledets (1968), standard dynamical [illegible]

1.2 Layer 2: TIG-Specific Claims with Strong Evidence

ID	Assumption	Statement	Status
A4	Quadratic iterate behavior partitions into exactly seven classes.	The discriminant (Delta), Lyapunov exponent (Lambda), and fixed-point stability divide all quadratic maps into [illegible]	Verified numerically across >10,000 operators (V[illegible]
A5	The PF dominant eigenfunction peaks at the stable fixed point.	For any quadratic map with |O'(x*)| < 1, the Perron-Frobenius invariant measure concentrates at x* [illegible]	Verified: 99.3% of 137 stable-fp operators in lattice-[illegible]
A6	The Lyapunov convergence formula matches iterate counts.	n = ceil[(ln|x0 - x*| + ln(1/epsilon)) / |Lambda|] reproduces the heuristic iteration count [illegible]	Verified: 85.0% within +/-2, 97.6% within +/-5 of 127[illegible]
A7	Routing along the invariant measure is optimal.	The coherence_router, which classifies path dynamics and routes along the PF eigenfunction, achieves [illegible]	Verified on synthetic benchmarks (Paper 2). Not yet [illegible]

1.3 Layer 3: TIG-Specific Claims Needing Further Evidence

ID	Assumption	Statement	Status
A8	The core equation S* = sigma(1-sig[illegible]	A [illegible] determined by coupling strength (sigma), self-ref[illegible]	Self-consisten[illegible]
A9	Multi-scale invariance holds.	The same equation produces S* > 0.8 from ecological to quantum scales [illegible]	[illegible]
A10	Ten operators form a complete basis.	The operators 0-9, with specific (a,b,c) coefficients, span all qualitative be[illegible]	F[illegible]

1.4 Layer 4: Forward-Looking Claims (Speculative)

ID	Assumption	Statement	Status
A11	The band classification generalizes beyond quadratic maps	Higher-order polynomial maps, coupled systems, and PDE discretizations should exhibit the same seven-band spectral class	Projected. Not yet tested. Would follow from A1 (ass
A12	Coherent intelligence is a distinct paradigm from statistical learning	Systems that organize by spectral classification of their dynamics (coherent intelligence) have	Argued in this paper (Section 3). No formal proof.

Peer review strategy: A reviewer who accepts A1-A3 (established math) and verifies A4-A7 (numerically testable) has the foundation. A8-A10 require independent reproduction. A11-A12 are the forward edge and should be evaluated as conjectures, not theorems.

2. What Follows from the Assumptions

We separate consequences into three tiers: proven (follows deductively from verified assumptions), projected (follows if unverified assumptions hold), and speculative (requires new work).

2.1 Proven Consequences (from A1-A7)

ID	Consequence	Derivation
C1	Any quadratic dynamical system has a rigorous spectral classification	From A1-A4. The seven bands are not heuristic categories; they are regions of the Perron-Frobenius spectrum, parti
C2	The natural density of a converging system is computable and peaked	From A2, A5. The PF eigenfunction gives the probability distribution over states. For stable systems, it is localized at
C3	Convergence rate is predictable from the spectrum	From A3, A6. Given |O'(x*)| = lambda, the number of steps to convergence is calculable. No simulation needed; the s
C4	Routing along the invariant measure outperforms blind algorithms	From A2, A7. A system that routes computation, traffic, or resources along the PF eigenfunction exploits the natural

2.2 Projected Consequences (require A8-A10)

ID	Consequence	What It Requires
C5	Multi-scale coherence monitoring is possible.	If A8-A9 hold, then a single scalar S* can monitor the coherence of any system at any scale. This would enable real-t
C6	System design can be optimized by band targeting.	If A10 holds (complete basis), then any desired system behavior can be constructed by composing operators from th
C7	Failure prediction becomes spectral gap monitoring.	If A8 holds, a declining spectral gap (S* trending toward T* = 0.714) signals impending transition from coherent to inc

2.3 Speculative Consequences (require A11-A12 and new work)

ID	Consequence	Argument
C8	Current AI architectures have a structural stability ceiling.	Neural networks optimize via gradient descent on a loss surface. This process has no spectral gap guarantee — the
C9	Coherent intelligence would be self-monitoring by construction.	A system organized by PF spectral classification knows its own stability state at every moment. The spectral gap IS t
C10	The alignment problem may be a spectral gap problem.	Alignment asks: how do we ensure a system's behavior remains within desired bounds? TIG reframes this: a syste

C8-C10 are the strongest claims in this paper and the most likely to be challenged. They are presented as conjectures with supporting arguments, not as proven theorems. The distinction matters.

3. Statistical Intelligence vs Coherent Intelligence

This section argues that TIG's results point toward a paradigm for intelligence that is structurally distinct from current artificial intelligence. We present this as an argument, not a proof.

3.1 What Current AI Does

Modern AI systems — large language models, diffusion models, reinforcement learning agents — share a common architecture: a parameterized function (neural network) is optimized by gradient descent to minimize a loss function over a dataset. The system learns statistical regularities in data. It does not model the dynamics of the system that generated the data.

This approach has produced extraordinary results. It has also produced characteristic failure modes:

Failure Mode	Mechanism
Hallucination	The model generates confident outputs that are factually wrong. This occurs because the model has learned statistical patterns, not causal
Alignment drift	A model trained on one objective may pursue that objective in unintended ways as its capability increases. The loss function constrains beh
Brittleness	Small changes to input distribution cause large changes in output quality. The model has no spectral gap guarantee — it does not know wh
Scaling uncertainty	Increasing model size improves average performance but does not guarantee stability. There is no theorem that says a larger model is mor

3.2 What Coherent Intelligence Would Do

A coherent intelligence system, as implied by TIG's framework, would differ from statistical AI in the following structural ways:

Property	Description
Organize by dynamics, not data	Instead of learning patterns from a dataset, the system would classify the spectral properties of its own state evolution. It wou
Route along invariant measures	Instead of forward-passing through layers, computation would flow along the PF eigenfunction — the natural density of the sy
Self-monitor via spectral gap	The system's stability metric (the PF spectral gap) is computable at every step. A declining gap is a measurable, advance wa
Structural alignment	Coherence above T* is a mathematical constraint, not a training objective. A system designed to maintain spectral gap > thre
Scale-invariant architecture	If A9 holds (multi-scale invariance), the same coherence equation governs behavior at every scale. There is no separate alig

3.3 The Distinction Is Structural, Not Incremental

We are not arguing that TIG is a better version of AI. We are arguing that coherent intelligence and statistical intelligence are different mathematical objects. Statistical AI optimizes a parameterized function over a data distribution. Coherent intelligence classifies and routes along the spectral structure of dynamics. The former has no stability theorem. The latter has one (the PF spectral gap). This is the core distinction.

An analogy: classical mechanics and quantum mechanics are both theories of motion, but they are structurally different theories. You cannot get quantum mechanics by making classical mechanics more accurate. Similarly, we conjecture that you cannot get coherent intelligence by making statistical AI larger or better-trained. The paradigm shift is from *learning patterns in data* to *classifying dynamics of systems*.

This is the speculative core of the paper. The analogy is suggestive, not proven. The concrete prediction is testable: build a system that routes computation along PF eigenfunctions and compare its stability properties to a neural network of equivalent capability. If the PF-routed system maintains coherence where the neural network does not, the paradigm distinction is real.

4. Implications (Tiered by Confidence)

4.1 Near-Term (achievable with current results)

1. Coherence-routed infrastructure: deploy the coherence_router on real networks. If the synthetic 100% throughput result holds in production, this is an immediate commercial application with no speculative assumptions.

2. Spectral stability monitoring: any system modeled as a quadratic iterate lattice can be monitored for spectral gap decline. This is a new class of early-warning system for infrastructure, financial, or biological systems.

3. AI safety research tool: the PF spectral decomposition can be applied to neural network training dynamics. Monitor the spectral gap of the loss landscape during training. If it collapses, the model is entering an unstable regime. This does not require replacing AI — it augments existing safety tooling.

4.2 Medium-Term (requires A8-A10 verification)

1. System design by spectral targeting: instead of engineering mechanisms and hoping they produce desired behavior, design the desired PF spectrum first and then construct the operator lattice that produces it. This inverts the engineering process.

2. Unified coherence metric across domains: a single S* score that compares the coherence of a neural network, an ecosystem, a power grid, and an economy on the same scale. If multi-scale invariance holds, this is a universal stability metric.

3. Coherence-native computing: processors that route computation along invariant measures rather than through fixed instruction pipelines. The architecture is dynamic, not static.

4.3 Long-Term (requires A11-A12 and paradigm validation)

1. Coherent intelligence systems: agents that organize by spectral classification of their own dynamics, self-monitor via spectral gap, and maintain structural alignment without external constraint. These would not be AI systems improved by TIG — they would be a different kind of system entirely.

2. Resolution of the alignment problem: if alignment can be reframed as spectral gap maintenance (C10), then a coherent system's safety is a mathematical invariant, not a training target. The system cannot misalign without a measurable spectral signature.

3. Post-statistical paradigm: intelligence research shifts from 'learn patterns from data' to 'classify and route dynamics.' This does not make current AI obsolete overnight. It provides a different foundation for the next generation of intelligent systems.

> *Tier 4.3 is where this paper is most likely to be accused of overreach. We include it because the implications follow logically from the assumptions, and because intellectual honesty requires stating where an argument leads even when the conclusion is large. These are not predictions — they are conditional consequences. If the assumptions hold, these implications follow. Whether the assumptions hold is an empirical question.*

5. What We Are Not Claiming

To prevent misinterpretation:

• TIG does not claim to replace AI. Current AI systems work. They solve real problems. TIG identifies structural limitations of the statistical paradigm and proposes an alternative foundation. Both can coexist.

• TIG does not claim to have solved the alignment problem. It conjectures (C10) that alignment may be reframed as spectral gap maintenance. This conjecture is testable but unproven.

• TIG does not claim its spectral indices are quantum numbers. They are spectral indices of a classical transfer operator. The structural analogy with quantum mechanics is precise but it is an analogy (Paper 5).

• TIG does not claim independent verification. All results are from a single research group (one human, AI co-development). Independent reproduction is the single highest priority for establishing credibility.

• TIG does not claim sigma = 0.991, T* = 0.714, or D* = 0.543 are derived from first principles. They are measured or equation-derived constants whose deeper origins are open questions.

• TIG does not claim to be a theory of everything. It is a framework for classifying and routing dynamical systems using spectral decomposition. The scope is specific even if the applications are broad.

6. Path to Peer Review

For a reviewer evaluating TIG, we recommend the following order of assessment:

Step	What to Verify	How	Expected Outcome
Step 1	Verify A4 (seven-band classification)	Run any quadratic map through the classifier. Confirm it falls into one of seven categories. This takes minutes and re[illegible]	Falsifiable if some state [illegible]
Step 2	Verify A5 (eigenfunction localization)	Build the discretized PF matrix for any stable-fp quadratic map. Compute the dominant eigenfunction. Compa[illegible]	Our result: 99.3%. Reviewer should get comparable peak[illegible]
Step 3	Verify A6 (convergence formula)	Compare n = ceil[(ln|x0-x*| + ln(1/epsilon)) / |Lambda|] with actual iteration counts across a range of parameters[illegible]	Our result: 85% within +/-2. Known deviation near bifurc[illegible]
Step 4	Verify A7 (routing advantage)	Run the coherence_router benchmark. Compare throughput against round-robin and consistent hashing under the sa[illegible]	Our result: 100% vs 10-12%. Synthetic benchmark; rea[illegible]
Step 5	Evaluate A8-A10 (core equation, multi-scale, basis completeness)	These require deeper engagement. We suggest starting with the ARCH stack simulation and checking S* > 0.8[illegible]	This is where independent reproduction matters most[illegible]
Step 6	Evaluate A11-A12 (speculative layer)	These are conjectures. Evaluate the arguments on their merits. The testable prediction is: build a PF-routed system a[illegible]	This is future work, not current claim[illegible]

7. Conclusion

TIG rests on twelve explicit assumptions, seven of which are either established mathematics or numerically verified. From these, four proven consequences follow deductively. Three projected consequences require verification of the self-consistency claims. Three speculative consequences point toward a paradigm we call coherent intelligence.

The strongest near-term result is practical: the coherence_router achieves categorically superior performance by routing along the Perron-Frobenius invariant measure. This works today, requires no speculative assumptions, and can be tested by anyone with a Python installation.

The strongest long-term implication is theoretical: if dynamical systems can be classified, monitored, and stabilized through their PF spectral decomposition, then intelligence systems built on this foundation would have structural stability guarantees that statistical AI does not. This is not a claim that AI is wrong. It is a claim that there exists a different mathematical foundation for intelligence, one where alignment and stability are properties of the spectrum rather than objectives of training.

Whether that claim survives contact with the broader research community is an open question. This paper provides the assumptions, the evidence, and the argument. We invite scrutiny.

References

[1] Koopman, B.O. (1931). 'Hamiltonian Systems and Transformation in Hilbert Space.' PNAS 17(5): 315-318.

[2] Lasota, A. & Mackey, M.C. (1994). *Chaos, Fractals, and Noise.* Springer.

[3] Baladi, V. (2000). *Positive Transfer Operators and Decay of Correlations.* World Scientific.

[4] Strogatz, S.H. (1994). *Nonlinear Dynamics and Chaos.* Westview Press.

[5] Devaney, R.L. (1989). *An Introduction to Chaotic Dynamical Systems.* Addison-Wesley.

[6] Oseledets, V.I. (1968). 'A Multiplicative Ergodic Theorem.' Trans. Moscow Math. Soc. 19: 197-231.

[7] Mezic, I. (2005). 'Spectral Properties of Dynamical Systems.' Nonlinear Dynamics 41: 309-325.

[8] Budisic, M., Mohr, R., Mezic, I. (2012). 'Applied Koopmanism.' Chaos 22(4): 047510.

[9] McLean, B. (2026). TIG White Papers 1-5. 7Site LLC. DOI: 10.5281/zenodo.